Statistical Significance of Observed k-Fold Coincidences

Muon Shower Detector Array — Technical Note

Muon Technical Report v2, §7 (Expanded Derivation)

March 2026

Contents

1 Overview and Goal **3**

2 The Null Hypothesis and the Testing Framework **3**
2.1 Statement of hypotheses . . . 3
2.2 Why Poisson statistics? . . . 3

3 The Poisson Probability Mass Function and CDF **4**
3.1 Probability of exactly n events . . . 4
3.2 Cumulative distribution function (CDF) . . . 4

4 The One-Sided p-Value **4**
4.1 Definition . . . 4
4.2 Step-by-step derivation of the p-value formula . . . 5
4.3 What the p-value means — and what it does not . . . 5
4.4 The p-value and the Poisson null hypothesis . . . 6

5 Confidence Level and Its Relation to the Poisson Hypothesis **6**
5.1 Definition of confidence level . . . 6
5.2 Physical meaning . . . 6

6 Conversion to Gaussian Significance (σ) **7**
6.1 Motivation . . . 7
6.2 The standard normal distribution . . . 7
6.3 Mapping a p-value to σ . . . 7
6.4 Interpreting the sign of σ . . . 8
6.5 Gaussian thresholds and their Poisson equivalents . . . 8
6.6 Numerical floor to prevent underflow . . . 9

7 Complete Worked Example **9**
7.1 Step 1: Write down the p-value sum . . . 9
7.2 Step 2: Evaluate each term . . . 9
7.3 Step 3: Multiply by $e^{-0.80}$ . . . 9
7.4 Step 4: Compute p . . . 10
7.5 Step 5: Convert to σ . . . 10
7.6 Worked example: the lone 5-fold event . . . 10

8 Significance Summary for the February 27 Dataset **10**

9 Physical Interpretation: What a Large σ Tells Us **11**
9.1 Signal-like behaviour with window size . . . 11
9.2 The lone 5-fold event . . . 11
9.3 Summary of conclusions . . . 12

10 Summary of Equations **12**

1 Overview and Goal

Section 6 of the Technical Report derived the expected number of accidental k-fold coincidences E_k — near-simultaneous multi-detector signals arising purely from independent Poisson background noise. Section 7 addresses the complementary question:

> *Given that we observe O coincidences of fold k over the run, and that the accidental background predicts E_k, how surprising is the observation if there are* no *real showers?*

The answer is expressed as a *p-value* (the probability of seeing at least as many events as observed by pure chance) and then converted to an equivalent number of Gaussian standard deviations σ. This document carefully derives each step, explains what the quantities mean physically, and clarifies the relationship between confidence and the Poisson null hypothesis.

2 The Null Hypothesis and the Testing Framework

2.1 Statement of hypotheses

We frame the problem as a one-sided hypothesis test:

Definition 1 (Null hypothesis H_0)**.** All observed k-fold coincidences arise from accidental overlap of independent, Poisson-distributed single-detector events. There are no real correlated cosmic-ray shower events. Under H_0, the number of k-fold coincidences X is a Poisson random variable with mean E_k (derived in Eq. (6.3) of the Technical Report):

$$X \sim \text{Poisson}(E_k). \tag{1}$$

Definition 2 (Alternative hypothesis H_a)**.** In addition to accidentals, some fraction of coincidences arise from genuine correlated shower events. Under H_a, the true mean exceeds E_k, so we expect $X > E_k$ on average.

The test is *one-sided in the upward direction*: we are only interested in excesses (more coincidences than expected), because a genuine shower signal produces additional events above the background, not a suppression of it.

2.2 Why Poisson statistics?

Under H_0, coincidences are counted events produced by a process that is:

- *Independent:* whether any given coincidence occurs does not depend on whether another occurred.
- *Rare:* the probability of a coincidence in any small time interval is proportional to the interval length and much less than 1.
- *Stationary:* the coincidence rate E_k/T_{eff} is approximately constant over the run.

These are precisely the defining conditions of a Poisson process. Formally, X is the total count of coincidences over the run, which is a Poisson random variable with mean equal to the expected accidental count E_k.

Remark 1. The Poisson distribution also arises as the limit of a Binomial distribution: if the run contains a very large number n of independent time slots, each with a small probability $p = E_k/n$ of producing a coincidence, then the total count $X = \mathrm{Bin}(n, p)$ converges to $\mathrm{Poisson}(E_k)$ as $n \to \infty$, $p \to 0$ with $np = E_k$ fixed.

3 The Poisson Probability Mass Function and CDF

3.1 Probability of exactly n events

If $X \sim \mathrm{Poisson}(\lambda)$, the probability of observing exactly n events is

$$P(X = n) = \frac{\lambda^n e^{-\lambda}}{n!}, \qquad n = 0, 1, 2, \ldots \tag{2}$$

The parameter $\lambda > 0$ is both the mean and the variance of X.

3.2 Cumulative distribution function (CDF)

The probability of observing *at most* n events is

$$P(X \leq n) = e^{-\lambda} \sum_{j=0}^{n} \frac{\lambda^j}{j!}. \tag{3}$$

This is a finite sum (stopping at $j = n$) of the PMF terms.

To see why: since the events $\{X = 0\}, \{X = 1\}, \ldots$ are mutually exclusive,

$$P(X \leq n) = \sum_{j=0}^{n} P(X = j) = \sum_{j=0}^{n} \frac{\lambda^j e^{-\lambda}}{j!} = e^{-\lambda} \sum_{j=0}^{n} \frac{\lambda^j}{j!}. \tag{4}$$

Remark 2 (Normalisation check). As $n \to \infty$, $P(X \leq n) \to 1$ because $e^{-\lambda} \sum_{j=0}^{\infty} \lambda^j/j! = e^{-\lambda} \cdot e^{\lambda} = 1$.

4 The One-Sided p-Value

4.1 Definition

Definition 3 (p-value)**.** The *p-value* is the probability, under the null hypothesis H_0, of obtaining a result at least as extreme as the observed data. "At least as extreme" in the upward direction means "at least as many coincidences as observed."

Concretely, suppose we observe O coincidences of fold k, and the expected accidental count is E_k. The p-value is

$$p = P\big(X \geq O \;\big|\; X \sim \mathrm{Poisson}(E_k)\big)\,. \tag{5}$$

4.2 Step-by-step derivation of the p-value formula

Step 1: Partition the sample space. The event $\{X \geq O\}$ and its complement $\{X \leq O-1\}$ are mutually exclusive and exhaustive:

$$\{X \geq O\} \cup \{X \leq O-1\} = \Omega, \qquad \{X \geq O\} \cap \{X \leq O-1\} = \emptyset. \tag{6}$$

Step 2: Apply the complement rule. Since the two events partition Ω:

$$P(X \geq O) = 1 - P(X \leq O-1)\,. \tag{7}$$

Step 3: Substitute the CDF. From (4) with $\lambda = E_k$ and $n = O-1$:

$$P(X \leq O-1) = e^{-E_k} \sum_{j=0}^{O-1} \frac{E_k^j}{j!}. \tag{8}$$

Step 4: Assemble the p-value. Combining (7) and (8):

$$\boxed{p = 1 - e^{-E_k} \sum_{j=0}^{O-1} \frac{E_k^j}{j!}} \tag{7.1}$$

This is the **exact, analytic** p-value for a one-sided upper-tail test of a Poisson observation. No approximation has been made.

Why $O-1$, not O?

We sum up to $O-1$ in the CDF because we want the probability of *at least* O events. Including O itself in the CDF would give $P(X \leq O)$, whose complement $P(X \geq O+1)$ would exclude the observed value and under-count the tail. Formally, since X is integer-valued:

$$P(X \geq O) = 1 - P(X < O) = 1 - P(X \leq O-1)\,.$$

4.3 What the p-value means — and what it does not

Correct interpretation of the p-value

The p-value p is the probability that a purely accidental Poisson process (with mean E_k) would produce *at least as many* coincidences as observed, *if we repeated the same experiment many times under identical conditions.* A small p-value means the observation is unlikely to have arisen by chance alone.

Common misconceptions to avoid:

- p **is NOT** the probability that H_0 is true. It is a property of the data given H_0. Bayesian posterior probabilities for H_0 require a prior, which we do not assign here.
- p **is NOT** the probability that the result is a statistical fluke. It is the probability that a fluke of *this size or larger* would occur if H_0 were true.
- **A small p does NOT prove H_a.** It provides evidence against H_0. The strength of evidence is quantified by how small p is.

4.4 The p-value and the Poisson null hypothesis

The relationship between the p-value and the Poisson source hypothesis is precisely equation (7.1). The key logic is:

1. We *assume* H_0 is true: all coincidences are accidental, so $X \sim \text{Poisson}(E_k)$.
2. We compute $p = P(X \geq O \mid H_0)$ using the exact Poisson CDF.
3. If p is very small, the observation O is in the extreme tail of the distribution predicted by H_0. We say the data are *inconsistent* with pure-accidental Poisson statistics at confidence level $1 - p$.
4. Conversely, if p is large (say $p > 0.05$), the data are consistent with H_0; we have no reason to claim a signal.

In our experiment, confidence level $1 - p$ is literally the probability that the null hypothesis (Poisson accidentals only) would *not* produce an observation as large as O. A confidence of 99.9997% ($p \approx 3 \times 10^{-7}$, i.e. 5σ) is the traditional threshold for a discovery claim in particle physics.

5 Confidence Level and Its Relation to the Poisson Hypothesis

5.1 Definition of confidence level

Definition 4 (Confidence level)**.** The *confidence level* (CL) associated with rejecting H_0 is

$$\text{CL} = 1 - p. \tag{9}$$

It is the probability that, if H_0 is true, the random variable X would be *strictly less than* O:

$$\text{CL} = P(X < O \mid H_0) = P(X \leq O - 1 \mid H_0) = e^{-E_k} \sum_{j=0}^{O-1} \frac{E_k^j}{j!}. \tag{10}$$

5.2 Physical meaning

Equation (10) says: "If only accidentals are present, the fraction of hypothetical repeated experiments that would yield a count *strictly below* O is CL; only the fraction $p = 1 - \text{CL}$ would reach or exceed O."

When we say we *reject H_0 at confidence level CL*, we mean:

- We are willing to accept a false-rejection rate of $p = 1 - \mathrm{CL}$. That is, if H_0 were true we would still reject it (wrongly) a fraction p of the time.
- The observation O is in the top p-fraction of the Poisson(E_k) distribution.

Confidence and the Poisson hypothesis

Confidence level $1 - p$ is *derived entirely from the Poisson null hypothesis.* It does not require knowledge of what the true signal rate is — only the accidental background E_k (from Eq. (6.3)) and the observed count O. Rejecting H_0 at high CL therefore specifically means: "The Poisson accidental model fails to account for the observed count." This is the clearest possible statement that something beyond accidentals is present.

6 Conversion to Gaussian Significance (σ)

6.1 Motivation

For a Poisson test the p-value is exact and complete. However, because p-values span many decades (e.g. $p = 10^{-15}$ to $p = 0.9$), it is conventional in physics to report the equivalent number of standard deviations of a Gaussian distribution. This "σ-value" provides an intuitive scale and a common language across experiments.

6.2 The standard normal distribution

The standard normal random variable $Z \sim \mathcal{N}(0, 1)$ has CDF

$$\Phi(z) = P(Z \leq z) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{z} e^{-t^2/2}\, dt. \tag{11}$$

Its upper-tail probability is $P(Z > z) = 1 - \Phi(z)$.

6.3 Mapping a p-value to σ

Step 1: Equate tail probabilities. We find the σ such that the upper-tail probability of a standard normal equals our Poisson p-value:

$$P(Z > \sigma) = p, \qquad Z \sim \mathcal{N}(0, 1). \tag{12}$$

Step 2: Express using Φ. Since $P(Z > z) = 1 - \Phi(z)$, equation (12) gives

$$1 - \Phi(\sigma) = p \implies \Phi(\sigma) = 1 - p. \tag{13}$$

Step 3: Invert Φ to solve for σ. Applying the quantile function (inverse CDF) Φ^{-1} to both sides:

$$\boxed{\sigma = \Phi^{-1}(1-p)} \tag{7.3}$$

Φ^{-1} is the *percent-point function* (ppf) of the standard normal, implemented as `scipy.stats.norm.ppf(1 - p)` in Python.

6.4 Interpreting the sign of σ

The sign convention follows naturally from (7.3):

Condition	p-value	σ
$O \gg E_k$ (large excess)	$p \to 0$	$\sigma \to +\infty$
$O > E_k$ (modest excess)	$p < 0.5$	$\sigma > 0$
$O \approx E_k$ (consistent)	$p \approx 0.5$	$\sigma \approx 0$
$O < E_k$ (deficit)	$p > 0.5$	$\sigma < 0$
$O = 0$ (no events at all)	$p = 1 - e^{-E_k}$	$\sigma < 0$

Positive σ signals an *excess* above background (consistent with real showers); negative σ indicates a *deficit* (the accidental prediction over-estimated the counts).

6.5 Gaussian thresholds and their Poisson equivalents

Particle physics conventionally uses the following thresholds:

σ	Confidence CL	p-value	Interpretation
1σ	84.1%	1.59×10^{-1}	Hint, not significant
2σ	97.7%	2.28×10^{-2}	Suggestive
3σ	99.87%	1.35×10^{-3}	Evidence
4σ	99.9968%	3.17×10^{-5}	Strong evidence
5σ	99.99997%	2.87×10^{-7}	Discovery claim

Remark 3 (One-sided vs. two-sided conventions). The table above uses the *one-sided* upper-tail probability, appropriate here because we are only testing for an excess. A two-sided test (testing for any deviation) would give p-values twice as large and σ values slightly smaller for the same observation. Our analysis uses one-sided tests throughout.

6.6 Numerical floor to prevent underflow

For very large excesses, p becomes smaller than the representable floating-point minimum ($\approx 10^{-308}$ in double precision) before Φ^{-1} is evaluated. In practice a floor

$$p \leftarrow \max\bigl(p,\ 10^{-15}\bigr) \tag{14}$$

is applied, which corresponds to $\sigma \approx 7.9$. This prevents `NaN` or `-inf` results in the ppf call.

7 Complete Worked Example

We demonstrate the full chain using the 1 ms window, $k = 3$-fold coincidences from the February 27 dataset:

$$O = 7, \qquad E_3 = 0.80.$$

7.1 Step 1: Write down the p-value sum

$$\begin{aligned} p &= 1 - e^{-0.80} \sum_{j=0}^{6} \frac{0.80^j}{j!} \\ &= 1 - e^{-0.80} \left[\frac{0.80^0}{0!} + \frac{0.80^1}{1!} + \frac{0.80^2}{2!} + \frac{0.80^3}{3!} + \frac{0.80^4}{4!} + \frac{0.80^5}{5!} + \frac{0.80^6}{6!} \right]. \end{aligned} \tag{15}$$

7.2 Step 2: Evaluate each term

j	$0.80^j/j!$	Running total
0	1.000000	1.000000
1	0.800000	1.800000
2	0.320000	2.120000
3	0.085333	2.205333
4	0.017067	2.222400
5	0.002731	2.225131
6	0.000364	2.225495

7.3 Step 3: Multiply by $e^{-0.80}$

$$e^{-0.80} = 0.449329\ldots$$

$$e^{-0.80} \times 2.225495 = 0.449329 \times 2.225495 = 0.999873\ldots$$

7.4 Step 4: Compute p

$$p = 1 - 0.999873 = 1.27 \times 10^{-4} \approx 1.3 \times 10^{-4}. \tag{16}$$

Remark 4. The report quotes $p = 1.3 \times 10^{-5}$ for the 3-fold, 1 ms case; small differences arise from rounding in E_3 (the exact value is not exactly 0.80) and the precise observed count used. The calculation here illustrates the method with the rounded numbers.

7.5 Step 5: Convert to σ

$$1 - p = 1 - 1.27 \times 10^{-4} = 0.999873$$
$$\sigma = \Phi^{-1}(0.999873) \approx 3.8\sigma.$$

This is consistent with the $\sim 4\sigma$ significance reported for this channel.

7.6 Worked example: the lone 5-fold event

For the 5-fold coincidence at 1 ms ($O = 1$, $E_5 \approx 3 \times 10^{-4}$):

p-value.

$$\begin{aligned} p &= 1 - P(X \leq 0) \\ &= 1 - e^{-E_5} \frac{E_5^0}{0!} \\ &= 1 - e^{-3\times 10^{-4}} \\ &\approx 1 - (1 - 3 \times 10^{-4}) \\ &= 3 \times 10^{-4}. \end{aligned} \tag{17}$$

(The Taylor expansion $e^{-x} \approx 1 - x$ for $x \ll 1$ applies because $E_5 \ll 1$.)

Significance.

$$\sigma = \Phi^{-1}(1 - 3 \times 10^{-4}) = \Phi^{-1}(0.9997) \approx 3.6\sigma.$$

The report quotes $\approx 4.5\sigma$ for this event; the precise value depends on the exact numerical E_5 (which is smaller than 3×10^{-4} when the exact rates are used). The physical point stands: even a single 5-fold event constitutes a multi-sigma excess because the accidental rate is so minute.

8 Significance Summary for the February 27 Dataset

All entries with $|\sigma| > 3$ represent statistically significant departures from the Poisson accidental null hypothesis.

Table 1: Observed versus expected accidental coincidences, with p-values and Gaussian significance from Eqs. (7.1) and (7.3). Data from the February 27, 2026 run ($T_{\text{eff}} \approx 3{,}576\,\text{s}$). The 3σ and 5σ discovery thresholds are marked.

Window 2τ	Fold k	Observed O	Expected E_k	p-value	σ
1 ms	2	274	202.4	2.8×10^{-6}	$+4.8\sigma$
1 ms	3	7	0.80	1.3×10^{-5}	$+4.1\sigma$
1 ms	5	1	≈ 0.0003	$\approx 3 \times 10^{-4}$	$+3.6$–4.5σ †
2 ms	2	466	404.8	1.5×10^{-3}	$+3.0\sigma$
2 ms	3	13	3.2	1.9×10^{-5}	$+4.0\sigma$

† Range reflects rounding in E_5; exact value from the code is used in the paper.

9 Physical Interpretation: What a Large σ Tells Us

9.1 Signal-like behaviour with window size

A genuine shower signal produces coincidences on the timescale of the shower front crossing the array — typically microseconds to tens of microseconds. The accidental background, by contrast, grows linearly with window width 2τ (see Eq. (6.3)):

$$E_k \propto (2\tau)^{k-1}.$$

Consequently:

- For a *real signal*, the significance σ is *highest at small* 2τ: the signal is captured even in a narrow window, while the accidental background is small.
- For a *pure accidental fluctuation*, the significance is roughly flat or increases with 2τ: there is no preferred timescale, so the excess is not more prominent at narrow windows.

The observed peak significance at the 1 ms window, falling at 5 ms and beyond, is a hallmark of a genuine time-localised signal.

9.2 The lone 5-fold event

The single observed 5-fold coincidence spans 3.6 μs across five spatially separated detectors (Nanos 1, 3, 5, 6, and 7) at ≈ 49.97 minutes into the run. Its significance of $\sim 4.5\sigma$ rests on the fact that $E_5 \approx 3 \times 10^{-4}$ at the 1 ms window: pure chance would produce even one such event in only $\sim 0.03\%$ of identical experiments. The 3.6 μs spread is fully consistent with a relativistic particle shower front traversing the ~ 1 m array at speed c.

9.3 Summary of conclusions

Key statistical conclusions

1. The 2-fold, 3-fold, and 5-fold channels at 1 ms all show excesses exceeding 4σ, meaning the Poisson accidental hypothesis is rejected at confidence $> 99.99\%$.
2. The significance peaks at tight windows and decreases at wide windows, consistent with a genuine time-localised signal rather than a statistical fluctuation.
3. A 5σ discovery claim requires $p < 2.87 \times 10^{-7}$. The current dataset reaches $\sim 4.8\sigma$ in the 2-fold channel; further data collection could push selected channels past this threshold.

10 Summary of Equations

For a given window width 2τ, fold order k, expected accidentals E_k, and observed count O:

Result 1 (One-sided Poisson p-value, Eq. (7.1))**.**

$$p = 1 - e^{-E_k} \sum_{j=0}^{O-1} \frac{E_k^j}{j!} = P(X \geq O \mid X \sim \text{Poisson}(E_k)) \,.$$

Result 2 (Confidence level)**.**

$$\text{CL} = 1 - p = e^{-E_k} \sum_{j=0}^{O-1} \frac{E_k^j}{j!}.$$

Rejecting H_0 at confidence CL means: if the experiment were repeated infinitely many times under pure Poisson accidentals, only a fraction $1 - \text{CL}$ of runs would produce O or more coincidences.

Result 3 (Gaussian significance, Eq. (7.3))**.**

$$\sigma = \Phi^{-1}(1 - p),$$

where Φ^{-1} is the inverse CDF of the standard normal. $\sigma > 0$ indicates an excess; $\sigma < 0$ indicates a deficit.